

Week 5: performance | QOL

**NRSC 7657 Workshop in Advanced Programming for
Neuroscientists**

course business

Performance

Processors: digital computers


vacuum tubes



ENIAC

18,000 vacuum tubes
5,000,000 hand-soldered joints
27 tons
150 kW

Transistors



Intel 4004

2250 transistors
MOS gates
< 27 tons
1W

Integrated circuits: CPUs



your computer, your phone

< 100,000 transistors
many Integrates gates
also < 27 tons
~100s of watts

Performance


Processors: CPU

- The CPU is (mostly) what is doing the heavy lifting when you ask python to do some calculation.
- The python interpreter converts your code to machine language (a set of instructions for what memory addresses to operate on), and then:
 1. Puts that machine language set of instruction in RAM
 2. the CPU does its thing (fetch-decode-execute-store)


Performance

Processors: CPU machine cycle

- Clock speed



Fetch-Decode-Execute Cycle (cont'd)



Performance


Processors: digital computers

vacuum tubes



ENIAC
100 kHz
Machine cycle: 200 μ s

Transistors



Intel 4004
740 kHz
Machine cycle: 10.8 μ s

Integrated circuits: CPUs



your computer, your phone
2 - 4Ghz
0.3 - 3 ns

Performance

Iteration: serial just like any other line

[2]: `a = 2`



[3]: `b = 2`

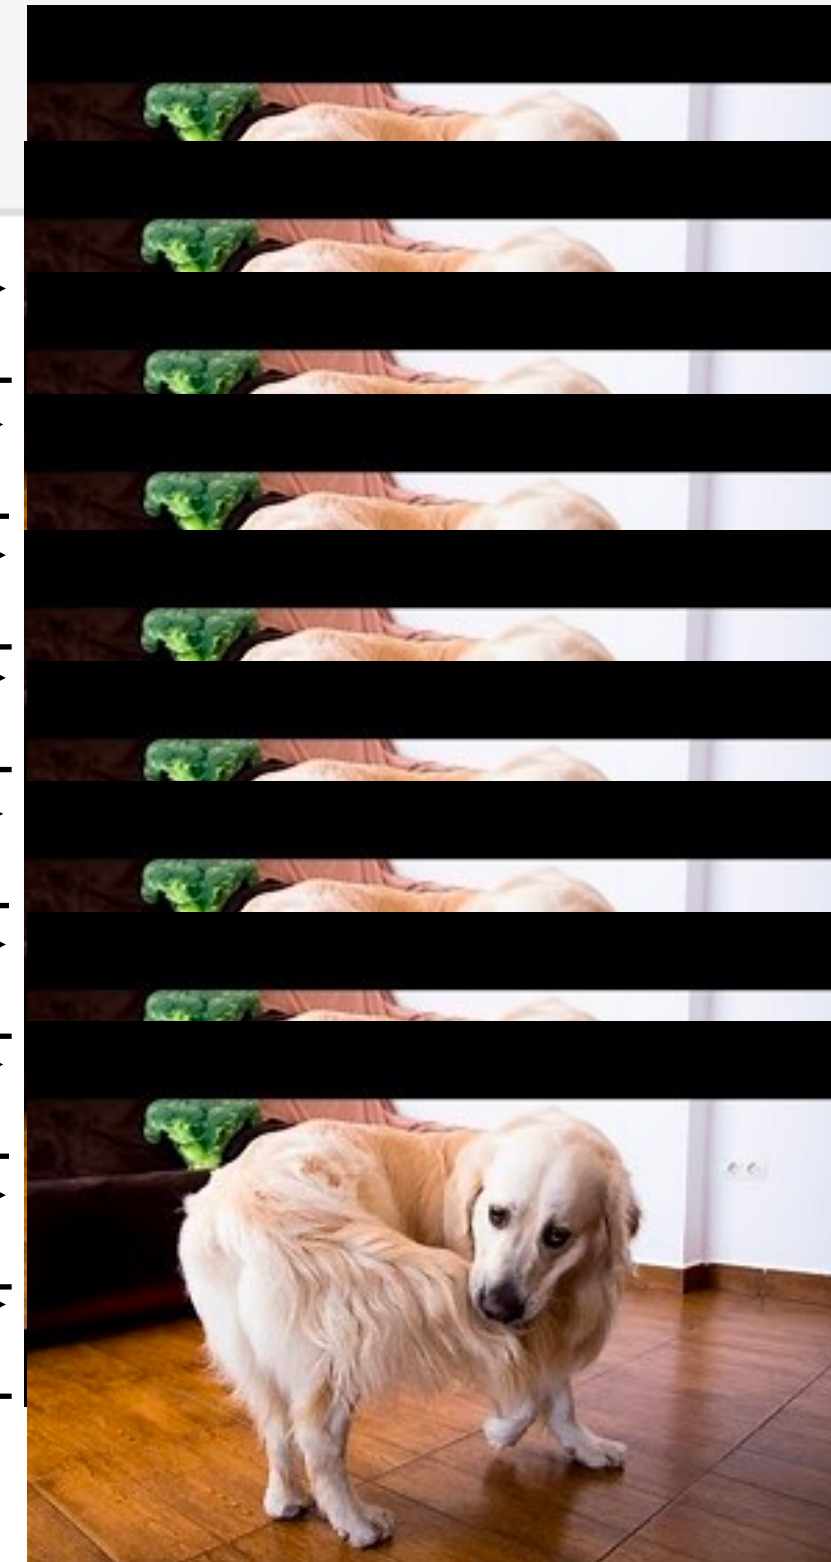



[4]: `c = a + b`

Performance

Iteration: serial just like any other line

```
[5]: for i in range(10):  
      print(i)
```



Performance

Iteration: serial just like any other line

```
[7]: print(0)
```

```
0
```

```
[8]: print(1)
```

```
1
```

```
[9]: print(2)
```

```
2
```

```
[10]: print(3)
```

```
3
```

```
[11]: print(4)
```

```
4
```

```
[12]: print(5)
```

```
5
```

```
[13]: print(6)
```

```
6
```

```
[14]: print(7)
```

```
7
```

```
[15]: print(8)
```

```
8
```

```
[16]: print(9)
```

```
9
```



Performance

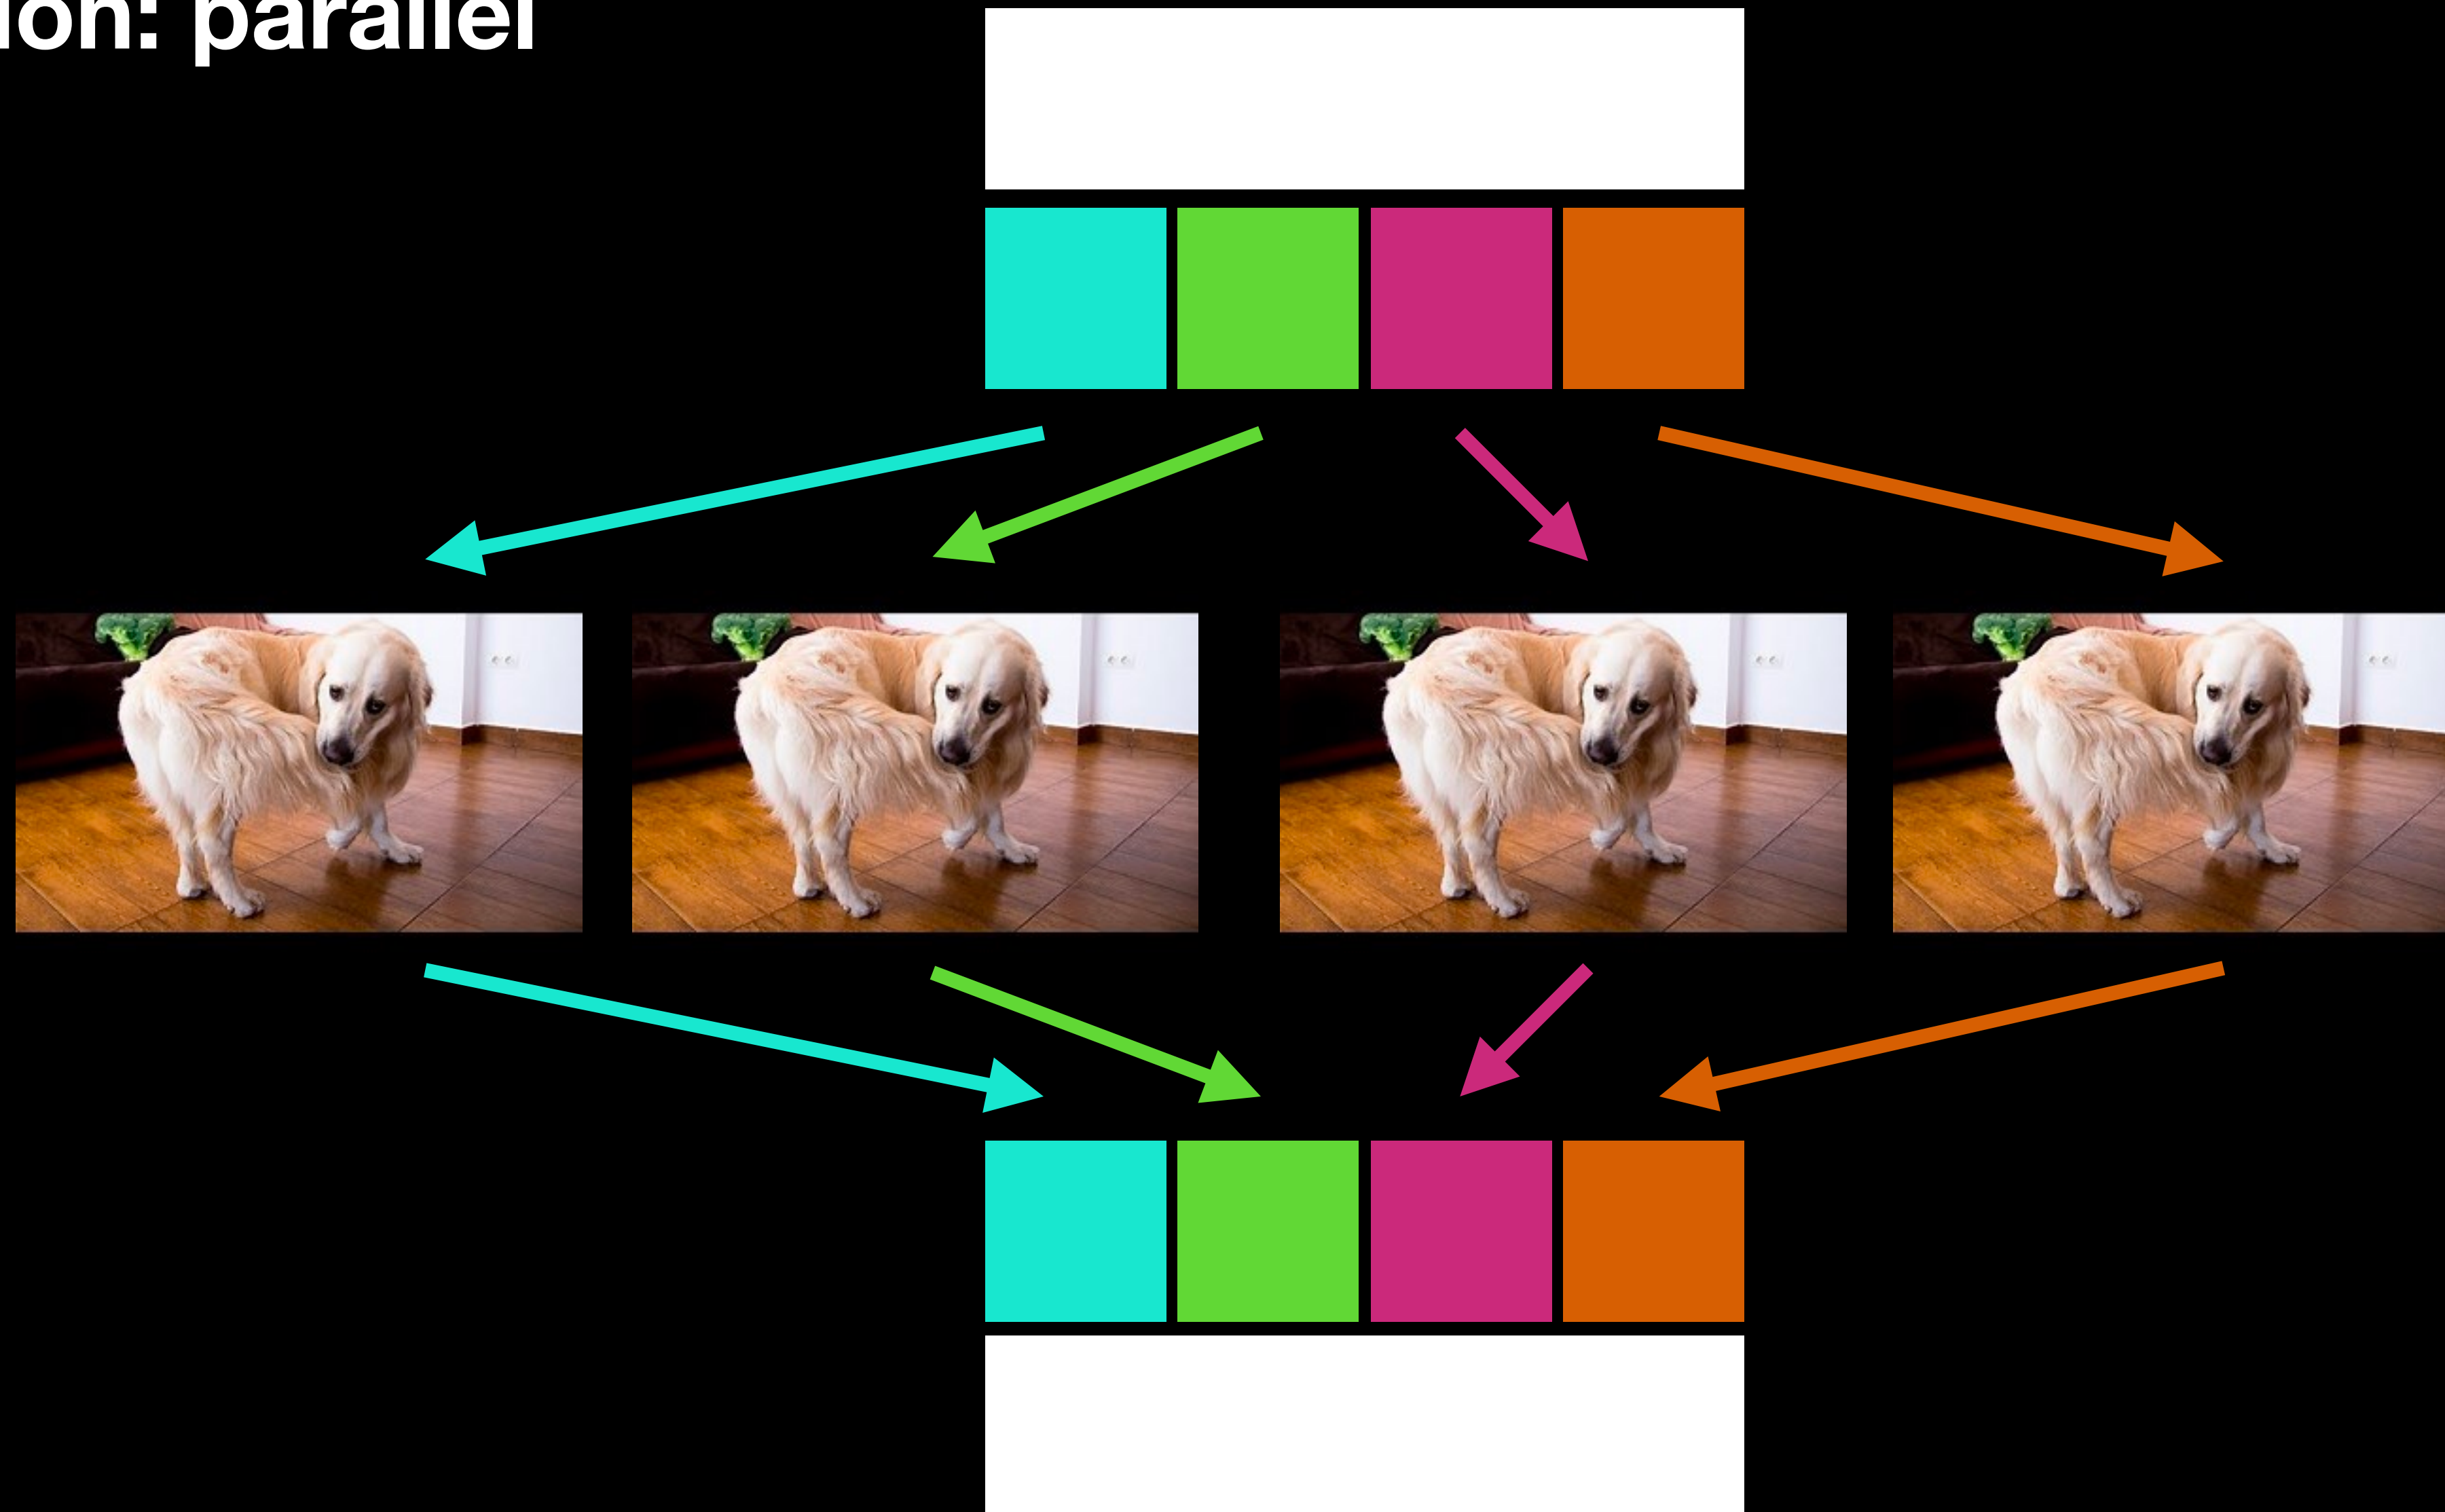
Iteration: data types, speeding it up

- Sorting tuples
- filtering
- Built-in methods (including filtering)
- List comprehension
- Avoiding it (but don't try too hard unless you need to)

[to be covered in .ipynb]

Performance

Iteration: parallel



Performance

Parallel CPU usage


- Split a computation into parts, ship it off to a number of processors, stitch the results back together
 - Not all computations benefit from this, but some do. The computations that be broken into parts that don't depend on each other are called **“embarrassingly parallel”**
 - Arrays where you do the same thing to chunks
 - Deep learning
- `multiprocessing`
 - `joblib`
 - `asyncio`
 - `dask`

Performance

Parallel CPU usage

With standard package helpers

Fetch-Decode-Execute Cycle (cont'd)




- multiprocessing
- joblib
- asyncio
- dask



Performance


Parallel CPU usage: dask

NumPy



Dask arrays scale NumPy workflows, enabling multi-dimensional data analysis in earth science, satellite imagery, genomics, biomedical applications, and machine learning algorithms.

pandas



Dask dataframes scale pandas workflows enabling applications in time series, business intelligence, and general data munging on big data.

```
# Arrays implement the NumPy API
import dask.array as da
x = da.random.random(size=(10000, 10000),
                      chunks=(1000, 1000))
x + x.T - x.mean(axis=0)
```

```
# Dataframes implement the pandas API
import dask.dataframe as dd
df = dd.read_csv('s3://.../2018-*-.csv')
df.groupby(df.account_id).balance.sum()
```

Google Cloud VMs

```
class dask_cloudprovider.gcp.GCPCluster(projectid=None, zone=None, network=None,
machine_type=None, on_host_maintenance=None, source_image=None, docker_image=None, ngpus=None,
gpu_type=None, filesystem_size=None, disk_type=None, auto_shutdown=None, bootstrap=True,
preemptible=None, debug=False, **kwargs) [source]
```

Cluster running on GCP VM Instances.

```
# Dask-ML implements the scikit-learn API
from dask_ml.linear_model \
    import LogisticRegression
lr = LogisticRegression()
lr.fit(train, test)
```



Performance

GPUs: many tiny cores for parallel things

The Difference between a CPU and GPU



CPU




GPU

- nVidia

Performance

GPUs: many tiny cores for parallel things



nVidia, <https://youtu.be/-P28LKWTzrl>

Performance

GPUs: many tiny cores for parallel things



- PyTorch
- TensorFlow
- MATLAB: Parallel Processing

Performance

Profiling

- `%%time` and `%%timeit`
 - Note: `%%timeit` will run multiple times by default, so use some caution

Code. QOL

Tips and tricks

- `tqdm`
- `os`, `sys`, `glob`
- `pprint`